

# Assignment 2: The N-Body Problem

Author: Nuno Nelas <fc51691@alunos.fc.ul.pt>

## Description of how parallelization was applied and rationale for the parallelization method used

Following the pattern of the previous assignment, this assignment focused a lot on a trial-and-error methodology, in which I've tried different parallelization methods to achieve the goal, which was to parallelize `NBodySystem.advance()` method.

Before even trying to parallelize it, there was a couple of things that I thought I could do to this method. On the assignment description we can find that this method has a  $O(n^2)$  complexity. However, after analyzing carefully this method, we can conclude that it has, in fact,  $O(n^2) + O(n)$ . This happens because there's a second for-loop to update x, y, z coordinates. To simplify this, I've extracted this logic to new method `void moveBody(NBody body, double dt)` and moved this inside the first for-loop.

I've also extracted velocity (vx, vz, vy) and mass calculations to a new method `void makeCalculations(NBody iBody, NBody body, double dt)`. As you can imagine, this minor refactor wasn't made to increase performance. Just to make things easier when trying different parallelization methods.

Regarding the parallelization itself, I will now describe the methods used the obtained results. (Whilst testing these different methods I've always compared them to the results achieved with the sequential version.)

- 1) Using bare-bone Threads: Since we were dealing with lists, I thought that we could do a somewhat similar approach to "Matrix Multiplication" exercise. However, I couldn't just split the for-loop into N different parts and hope this would work. Every iteration has to have a reference to its past elements, so as I should have expected, this led to inaccurate **energy** results.
- 2) Using ExecutorService: `ExecutorService` abstracts away many of the complexities associated with the lower-level abstractions like raw `Thread`. This seemed like a good choice. Basically, I've created a `newFixedThreadPool` with 8 number of threads (to match my CPU specs) and instead of calling `makeCalculations` directly, I called it with `executor.submit`. This was terrible performance wise, even to the point of stop running the simulation because it was running for too long with my CPU at 100% on all cores.
- 3) Using Streams: With Streams I found that it was somewhat simple to parallelize the `advance` method. In practice, I would only need to 1) call `.parallel()` to iterate the for-loop in parallel; 2) call `.forEachOrdered()` to ensure that the items were executed orderly.

## Measurements showing whether parallelization was advantageous in each case

For this assignment I have used a machine with the following configurations:

CPU	CPU Cores	CPU Threads	RAM	Operating System
Intel i7-8550U (8) @ 4.000GHz	4	8	24 GB	Linux Mint 20 x86_64

As the above screenshots from VisualVM show, although using Streams API simplifies all that "hard-work" of managing tasks, it does not resolve all synchronization issues nor makes the solution faster. A parallel stream should not be used when we have a massive amount of items to process, that only increases the thread coordination problem. To sum up, using Streams led to an increase of more than 1 minute between sequential and parallelized version.

